

Student Name <- replace with your name

## CS585 Spring 2023 Programming Assignment #01

Due: Sunday, February 11, 2024, 11:59 PM CST

Points: 150

### Instructions:

1. Place **all your deliverables (as described below) into a single ZIP** file named:

LastName\_FirstName\_CS585\_Programming01.zip

2. Submit it to Blackboard Assignments section before the due date. **No late submissions will be accepted.**

### Objectives:

1. (50 points) Perform basic word frequency distribution analysis for a text corpus.
2. (50 points) Calculate probability of a sentence.
3. (50 points) Language Model word prediction.

### Deliverables:

Your submission should include:

- **Make sure your code is sufficiently commented! NO Jupyter Notebook files!**
- **Part A:** Python code file(s). Your py file should be named:

cs585\_P01A\_XXXXXXXXX.py

where XXXXXXXXX is your IIT A number (**this is REQUIRED!**). If your solution uses multiple files, makes sure that the main (the one that will be run to solve the problem) is named that way and others include your IIT A number in their names as well.

- **Part B:** Python code file(s). Your py file should be named:

cs585\_P01B\_XXXXXXXXX.py

where XXXXXXXXX is your IIT A number (**this is REQUIRED!**). If your solution uses multiple files, makes sure that the main (the one that will be run to solve the problem) is named that way and others include your IIT A number in their names as well.

- **Part C:** Python code file(s). Your py file should be named:

cs585\_P01C\_XXXXXXXXX.py

where XXXXXXXXX is your IIT A number (**this is REQUIRED!**). If your solution uses multiple files, makes sure that the main (the one that will be run to solve the problem) is named that way and others include your IIT A number in their names as well.

- this document with your results and conclusions. You should rename it to:

LastName\_FirstName\_cs585\_Programming01.doc or pdf

MS WORD or PDF formats only, please.

### Part A [50 pts]:

Use Python's NLTK package along with the corpora:

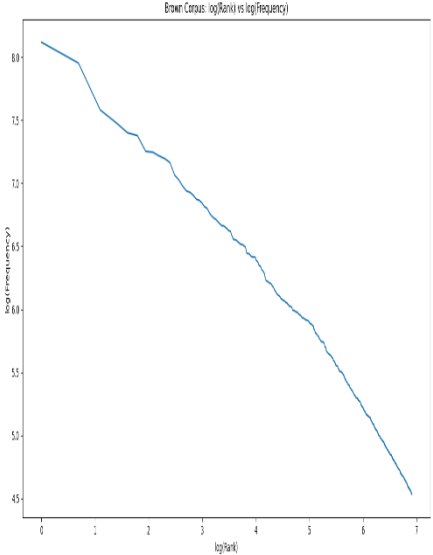
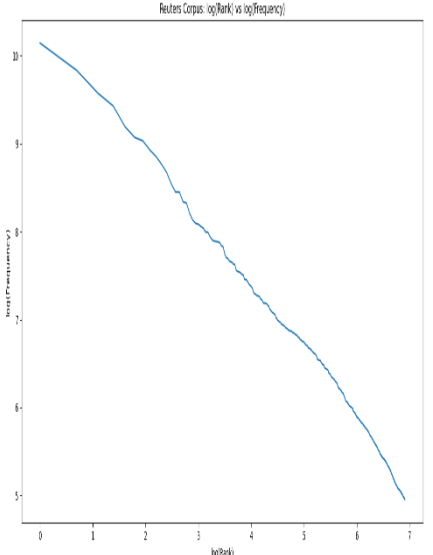
- Brown,
- Reuters,

to:

- 1) [10 pts] obtain the word frequency distribution (after removing all stop words; use the **stopwords** corpora for that purpose) for BOTH corpora,
- 2) [10 pts] display a top ten (ranks 1 through 10) words for BOTH corpora on screen (also place them in the table below)

Top 10 words					
Brown			Reuters		
Top Ten Words in the Brown Corpus:			Top Ten Words in the Reuters Corpus:		
Rank	Word	Frequency	Rank	Word	Frequency
1	one	3357	1	said	25383
2	would	2843	2	mln	18623
3	said	1961	3	vs	14341
4	could	1777	4	dlrs	12417
5	new	1635	5	pct	9810
6	time	1600	6	lt	8696
7	two	1412	7	cts	8361
8	may	1402	8	year	7529
9	first	1361	9	net	6989
10	man	1332	10	u	6392

3) [15 pts] generate **log(rank) vs log(frequency)** plots for the first 1000 (ranks 1 through 1000) words for BOTH corpora (you can use the matplotlib package or some other plotting package / tool). Place BOTH plots in the table below.

log(rank) vs log(frequency) plots	
Brown	Reuters
	
<p><b>Did you observe anything interesting when comparing all plots? Write your comments below:</b></p> <p>It has been noted that the Reuters Corpus has higher frequencies than the Brown Corpus, as the maximum value of log (Frequency) is 10, compared to 8 in the Brown Corpus.</p>	

4) [15 pts] use frequency counts obtained earlier to calculate the unigram occurrence probability for the TWO (“technical” and not technical) words. Use lowercasing first! **Display all relevant counts and probability on screen for BOTH corpora (also: enter final values in the table below)**. It can be zero for some words.

“technical” / seldom used in casual conversation word (for example “adiabatic”	
Brown	Reuters

Count: <b>120</b> Probability: <b>0.000233</b>	Count: <b>95</b> Probability: <b>0.000109</b>
<b>Non- technical / casual / daily-use word (for example “dinner”)</b>	
<b>Brown</b>	<b>Reuters</b>
Count: <b>0</b> Probability: <b>0.000000</b>	Count: <b>0</b> Probability: <b>0.000000</b>

### Part B [50 pts]:

Use Python’s NLTK package along with the Brown corpus for the following tasks:

1. **[1 pts]** Ask the user to enter a sentence S from a keyboard.
2. **[1 pts]** Apply lowercasing to S.
3. **[45 pts]** Calculate  $P(S)$  assuming a 2-gram language model (**assume that probability of any bigram starting or ending a sentence is 0.25**)
4. **[3 pts]** Display the sentence S, list all the individual bigrams and their probabilities, and the final probability  $P(S)$  on screen. It is fine if it is zero.

### Part C [50 pts]:

Use Python’s NLTK package along with the Brown corpus (after removing all stop words; use the **stopwords** corpora for that purpose) for the following tasks:

1. **[1 pts]** Start by asking the user for initial word/token W1. Apply lowercasing to W1 (and all future entries). If the word is NOT in the corpus offer two options:
  - a. **ask again**
  - b. **QUIT**
2. **[45 pts]** Assuming a 2-gram language model, a menu with TOP 3 “most likely to follow W1” words (according to the W1, NEXT WORD probability estimate). For example, if the user started with W1 = “good”, the following could be displayed (**NOTE: I made up this selection and corresponding probability estimates**):

good ...

Which word should follow:

- 1) morning  $P(\text{good morning}) = 0.25$
- 2) evening  $P(\text{good evening}) = 0.15$
- 3) afternoon  $P(\text{good afternoon}) = 0.14$
- 4) QUIT

**Repeat (and add subsequent word choices to the “sentence”) until user selects (4) and QUITs.**

If the user picks a number other than 1,2,3, and 4, **assume user choice is (1).**